

# Lecture 1 - R Basics

Make your paper figures professionally: Scientific data analysis and visualization with R

Julien Gagneur

- 1 Introduction to R and RStudio
- 2 First steps with R
- 3 Data Types in R
- 4 Sorting and ranking
- 5 Installing and loading packages
- 6 Data wrangling
- 7 Introduction to Data.tables
- 8 Row subsetting
- 9 Column operations

10 Grouping

11 Counting occurrences with `.N`

12 Extending tables

13 Rmarkdown

14 Recap

15 Reading

16 Cheatsheets

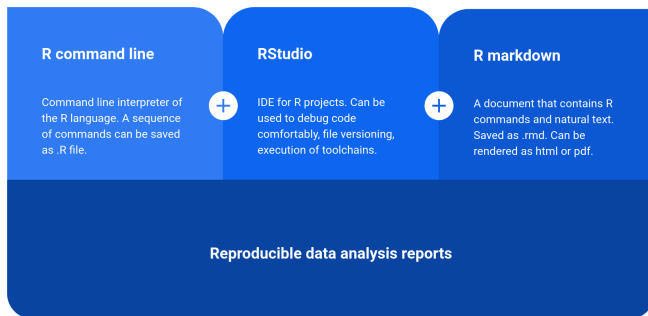
# Introduction to R and RStudio

# The programming language R

- R is an open source implementation of S (S-Plus is a commercial implementation)
- R is available under GNU Copy-left
- R is group project run by a core group of developers (with new releases semi-annually)

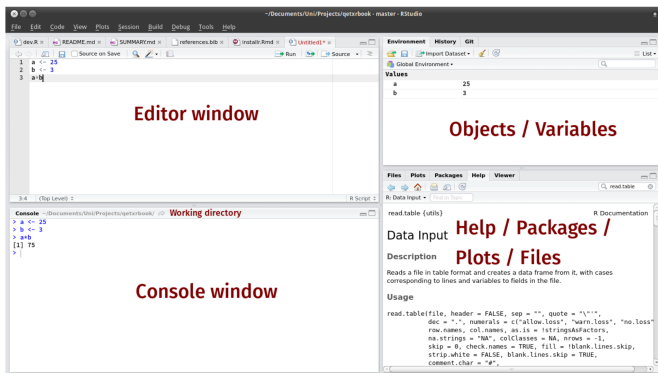
# R and RStudio

R markdown builds on top of R and RStudio



# Rstudio

- Rstudio is a software that allows to program in R and interactively analyse data with R
- It organizes the R session into 4 panes:



# First steps with R



## Disclaimer

This section is largely borrowed from the book Introduction to Data Science by Rafael Irizarry.  
[<https://rafalab.github.io/dsbook>]

You can find the whole book of our Data Analysis and Visualization in R lecture here:  
[<https://gagneurlab.github.io/dataviz/>]

A cheatsheet for R studio can be obtained here:  
[<https://raw.githubusercontent.com/rstudio/cheatsheets/master/rstudio-ide.pdf>]

While a cheatsheet for R basics can be obtained here:  
[<https://www.rstudio.com/wp-content/uploads/2016/10/r-cheat-sheet-3.pdf>]

# Assignments

All big (programming) journeys start with a small step (or assignment).

We use `<-` to assign **values** to **variables**.

```
a <- 9  
b <- 3 + 2
```

We can also assign values using `=` instead of `<-`, but we recommend against using `=` to avoid confusion.

# Objects

To see the value stored in a variable, we simply ask R to evaluate `a` and it shows the stored value:

```
a
```

```
## [1] 9
```

A more explicit way to ask R to show us the value stored in `a` is using `print` like this:

```
print(a)
```

```
## [1] 9
```

We use the term *object* to describe stuff that is stored in R. Variables are examples, but objects can also be more complicated entities such as functions.

# Functions

The data analysis process can usually be described as a series of *functions* applied to the data. R includes several predefined functions and most of the analysis pipelines we construct make extensive use of these.

For example, we can compute the square root of `a` with `sqrt` or see all the variables saved in our workspace by calling the function `ls`:

```
sqrt(a)
```

```
## [1] 3
```

```
ls()
```

```
## [1] "a" "b"
```

Unlike `ls`, most functions require one or more *arguments*.

In general, we need to use parentheses to evaluate a function. Without them, the function is not evaluated and instead R shows the code that defines the function:

```
sqrt
```

```
## function (x) .Primitive("sqrt")
```

We can find out what the function expects and what it does by reviewing the manuals included in R with the help of the shorthand ? (available for most functions):

```
?log
```

The help page will show us that `log` needs `x` and `base` to run and that the argument `base` is optional.

We can change the default values of optional arguments by simply assigning another object:

```
log(8, base = 2)
```

If no argument name is used, R assumes we are entering arguments in the order shown in the help file. So by not using the names, it assumes the arguments are `x` followed by `base`:

```
log(8,2)
```

If using the arguments' names, then we can include them in whatever order we want:

```
log(base = 2, x = 8)
```

To specify arguments, we must use `=`, and cannot use `<-`.

# Data Types in R

# Data types

Variables in R can be of different types. For example, we need to distinguish numbers from character strings and tables from simple lists of numbers. The function `class` helps us determine what type of object we have:

```
a <- 2  
class(a)
```

```
## [1] "numeric"
```

To work efficiently in R, it is important to learn the different types of variables and what we can do with these.



## Vectors: numerics, characters, and logical

```
my_vector
```

```
## [1] 1 2 3 4 5
```

The object `my_vector` is not one number but several. We call these types of objects *vectors*, which can be stored as variables.

The function `length` tells you how many entries are in the vector:

```
length(my_vector)
```

```
## [1] 5
```

This particular vector is *numeric* since it contains numbers:

```
class(my_vector)
```

```
## [1] "numeric"
```

To store character strings, vectors can also be of class *character*. For example, we can create a vector containing strings as follows:

```
char_vec <- c('DatViz', 'is', 'cool')  
char_vec
```

```
## [1] "DatViz" "is"      "cool"
```

```
class(char_vec)
```

```
## [1] "character"
```

Note that we can use the function `c()`, which stands for *concatenate*, to create vectors of any type.

Another important type of vectors are *logical vectors*. These must be either TRUE or FALSE.

```
z <- c(3 == 2, 5>4)
```

```
z
```

```
## [1] FALSE  TRUE
```

```
class(z)
```

```
## [1] "logical"
```

The == is a relational operator asking if 3 is equal to 2.

If we just use one =, we actually assign a variable, but if we use two == you test for equality.

## Naming vectors

Sometimes it is useful to name the entries of a vector.

For example, when defining a vector of country codes, we can use the names to connect the two:

```
codes <- c(italy = 380, canada = 124, egypt = 818)
codes
```

```
##  italy canada  egypt
##    380    124    818
```

We can also assign names to an unnamed vector using the `names` function:

```
codes <- c(380, 124, 818)
country <- c("italy", "canada", "egypt")
names(codes) <- country
```

## Vectors of sequences

Another useful function for creating vectors generates sequences:

```
seq(1, 10)
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

The first argument defines the start, and the second defines the end which is included. The default is to go up in increments of 1, but a third argument lets us tell it how much to jump by:

```
seq(1, 10, 2)
```

```
## [1] 1 3 5 7 9
```

## Vectors containing repetitions

The `rep` function replicates values for a specific number of times. It can be useful when want to create a vector that contains repetitions.

For example, we can create the following vector with the `c` function:

```
x <- c(1,2,3,1,2,3,1,2,3,1,2,3)
```

But we can also create the same vector much easier with `rep`:

```
x <- rep(1:3, times=4)
```

```
x
```

```
## [1] 1 2 3 1 2 3 1 2 3 1 2 3
```

We can also pass a vector to the `rep` function and tell it that we want each entry to be repeated a certain number of times:

```
s <- c("Jump", 'Go')  
x <- rep(s, each=3)  
x
```

```
## [1] "Jump" "Jump" "Jump" "Go"   "Go"   "Go"
```

Or we can define the output length and let R figure out how many times it should repeat the entries in the given vector:

```
x <- rep(c(TRUE,FALSE,FALSE), len=10)  
x
```

```
## [1] TRUE FALSE FALSE TRUE FALSE FALSE TRUE FALSE FALSE TRUE
```

## Subsetting vectors

We use square brackets to access specific elements of a vector.

For instance, we can access the second element of a vector using:

```
codes[2]
```

```
## canada  
##      124
```

We can get more than one entry by using a multi-entry vector as an index:

```
codes[c(1,3)]
```

```
## italy egypt  
##   380   818
```



We can access consecutive entries in a vector:

```
codes[1:2]
```

```
##  italy canada  
##    380    124
```

If the elements have names, we can also access the entries using these names. Below are two examples.

```
codes["canada"]
```

```
##  canada  
##    124
```

```
codes[c("egypt", "italy")]
```

```
##  egypt italy  
##    818    380
```

## Rescaling a vector

In R, arithmetic operations on vectors occur *element-wise*.

For a quick example, we can convert a vector containing height values in inches to centimeters:

```
inches <- c(69, 62, 66, 70, 70, 73, 67, 73, 67, 70)
inches * 2.54
```

```
## [1] 175.26 157.48 167.64 177.80 177.80 185.42 170.18 185.42 170.18 177.80
```

We can not also multiply a vector times a scalar, but also perform additions and subtractions:

```
inches - 69
```

```
## [1] 0 -7 -3 1 1 4 -2 4 -2 1
```

## Arithmetics with two vectors

If we have two vectors of the same length, and we sum them in R, they will be added entry by entry as follows:

$$\begin{pmatrix} a \\ b \\ c \\ d \end{pmatrix} + \begin{pmatrix} e \\ f \\ g \\ h \end{pmatrix} = \begin{pmatrix} a + e \\ b + f \\ c + g \\ d + h \end{pmatrix}$$

The same holds for other mathematical operations, such as  $-$ ,  $*$  and  $/$ .

# Coercion in R

In general, *coercion* is an attempt by R to be flexible with data types.

When an entry does not match the expected, some of the prebuilt R functions try to guess what was meant before throwing an error.

We said that vectors must be all of the same type. So if we try to combine, say, numbers and characters, you might expect an error:

```
x <- c(1, "canada", 3)
```

But we don't get one, not even a warning! What happened? Look at x and its class:

```
x  
## [1] "1"      "canada" "3"  
class(x)
```

```
## [1] "character"
```

R *coerced* the data into characters!

## Not availables (NA)

When a function tries to coerce one type to another and encounters an impossible case, it usually gives us a warning and turns the entry into a special value called an NA for “not available”. For example:

```
x <- c("1", "b", "3")  
as.numeric(x)
```

```
## Warning: NAs introduced by coercion
```

```
## [1] 1 NA 3
```

R does not have any guesses for what number we want when you type b, so it does not try.

We will encounter the NAs often as they are generally used for missing data, a common problem in real-world datasets.

# Vector exercises

# Factors

```
my_factor
```

```
## [1] Mutant Mutant Mutant WT      WT      WT  
## Levels: Mutant WT
```

The `my_factor` variable, might look like a character vector. However, it is a *factor*:

```
class(my_factor)
```

```
## [1] "factor"
```

Factors are useful for storing categorical data.

We can inspect the categories (or levels) of a factor by using the `levels` function:

```
levels(my_factor)
```

```
## [1] "Mutant" "WT"
```



By default the levels are the unique values, sorted by alphanumerical order. We can construct a factor as follows:

```
dogs <- factor(c('Beagle', 'Poodle',  
                'Labrador', 'Beagle', 'Akita'))  
dogs
```

```
## [1] Beagle  Poodle  Labrador Beagle  Akita  
## Levels: Akita Beagle Labrador Poodle
```

In the background, R stores these *levels* as integers and keeps a map to keep track of the labels. This is more memory efficient than storing all the characters.

## Further data types

Other data types in R include:

- **lists** as the generalization of data frames
- **matrices** for two dimensional data

See the script for more information about them!

# Factor exercises



## Sorting and ranking

# Sorting

The function `sort` sorts a vector in increasing order:

```
my_vector <- c(6, 1, 2, 5, 10, 9, 8)
sort(my_vector)
```

```
## [1] 1 2 5 6 8 9 10
```

or in decreasing order:

```
sort(my_vector, decreasing = TRUE)
```

```
## [1] 10 9 8 6 5 2 1
```

# Ordering

The function `order` takes a vector as input and, rather than sorting the input vector, it returns the index that sorts input vector:

```
x <- c(31, 4, 15, 92, 65)
index <- order(x)
index
```

```
## [1] 2 3 1 5 4
x[index]
```

```
## [1] 4 15 31 65 92
```

This is the same output as that returned by `sort(x)`.

## max and which.max

If we are only interested in the entry with the largest value, we can use `max` for the value:

```
my_vector <- c(6, 1, 2, 5, 10, 9, 8)
max(my_vector)
```

```
## [1] 10
```

and `which.max` for the index of the largest value:

```
i_max <- which.max(my_vector)
i_max
```

```
## [1] 5
```

```
my_vector[i_max]
```

```
## [1] 10
```

For the minimum, we can use `min` and `which.min` in the same way.



# Ranking

The function `rank` is also related to order and can be useful. For any given vector it returns a vector with the rank of the first entry, second entry, etc., of the input vector. Here is a simple example:

```
x <- c(31, 4, 15, 92, 65)  
rank(x)
```

```
## [1] 3 1 2 5 4
```

## Installing and loading packages

# Installing and loading packages

Packages are the fundamental units of reproducible R code. Several packages are automatically included.

We can install and load new packages by typing:

```
install.packages("vegan") # install new package called vegan  
library(vegan) # and load it
```

Vegan is a package to analyze biodiversity. To learn more about an installed package try:

```
browseVignettes("vegan")
```

## Curious about learning more R basics?

- Read the first chapter and appendix of our script!
- Ask questions on Slack!
- Practice with DataCamp!

# Data wrangling

# Data wrangling

- Data wrangling refers to the task of processing raw data into useful formats
- This Chapter introduces basic data wrangling operations in R using `data.table` from the R package `data.table`:

```
# install.packages("data.table")  
library(data.table)
```

A cheatsheet for R basics can be obtained here:

[<https://www.rstudio.com/wp-content/uploads/2016/10/r-cheat-sheet-3.pdf>]

A cheatsheet for simple `data.table` manipulations can be obtained here: [<https://datacamp-community-prod.s3.amazonaws.com/6fdf799f-76ba-45b1-b8d8-39c4d4211c31>]

# Introduction to Data.tables

# Overview

- `data.table` objects are a modern implementation of tables containing
  - variables stored in columns and
  - observations stored in rows
- A `data.table` is a memory efficient and faster implementation of `data.frame`.
  - more efficient because it operates on its columns by reference (without copying)
  - from now on: work only with `data.table`
- Each column can have a different type
- A `data.table` does not have row names
- Shorter and more flexible syntax than `data.frame`



## Basic data.table syntax

The general basic form of the data.table syntax is:

```
DT[ i, j, by ] #
  |  |  |
  |  |  -----> grouped by what?
  |  -----> what to do with the columns?
  ---> on which rows?
```

“Take DT, subset rows by i, then compute j grouped by by”.

## Creating data.tables

To create a `data.table`, we just name its columns and populate them:

```
library(data.table)
DT <- data.table(x = rep(c("a","b","c"), each = 3), y = c(1, 3, 6), v = 1:9)
DT # note how column y was recycled
```

```
##      x y v
## 1: a 1 1
## 2: a 3 2
## 3: a 6 3
## 4: b 1 4
## 5: b 3 5
## 6: b 6 6
## 7: c 1 7
## 8: c 3 8
## 9: c 6 9
```

All the columns have to have the same length.

If vectors of different lengths are provided upon creation of a `data.table`, R automatically recycles the values of the shorter vectors.

## Converting into data.table

If we want to convert any other R object to a data.table, all we have to do is to call the `as.data.table()` function.

This is typically done for data.frame objects:

```
#install.packages("dslabs")  
library(dslabs)  
brexit_polls <- as.data.table(brexit_polls)  
class(brexit_polls)
```

```
## [1] "data.table" "data.frame"
```

titanic\_df is now both a data.table and a data.frame as data.table inherits from data.frame

# Loading data.tables

- We can read files from disk and process them using `data.table`
- The easiest way to do so is to use the function `fread()`

## Loading data.tables

Example: Kaggle flight and airports dataset that is limited to flights going and in or to the Los Angeles airport:

```
flights <- fread('path_to_file/flightsLAX.csv')
```

```
head(flights, n=5)
```

##	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	TAIL_NUMBER	ORIGIN_AIRPORT
## 1:	2015	1	1	4	AA	2336	N3KUAA	LAX
## 2:	2015	1	1	4	AA	258	N3HYAA	LAX
## 3:	2015	1	1	4	US	2013	N584UW	LAX
## 4:	2015	1	1	4	DL	1434	N547US	LAX
## 5:	2015	1	1	4	AA	115	N3CTAA	LAX
##	DESTINATION_AIRPORT			DEPARTURE_TIME	AIR_TIME	DISTANCE	ARRIVAL_TIME	
## 1:	PBI			2	263	2330	741	
## 2:	MIA			15	258	2342	756	
## 3:	CLT			44	228	2125	753	
## 4:	MSP			35	188	1535	605	
## 5:	MIA			103	255	2342	839	

## Inspecting tables

A first step in any analysis should involve inspecting the data we just read in.

After looking at the first and last rows of the table, the next information we are often interested in is the **size** of our data set:

```
ncol(flights)    # nrow(flights) for number of rows
```

```
## [1] 13
```

```
dim(flights)     # returns nrow and ncol
```

```
## [1] 389369    13
```

## Basic statistics

Next, we are often interested in **basic statistics** on the columns.

To obtain this information we can call the `summary()` function on the table:

```
summary(flights[,1:6])
```

```
##           YEAR           MONTH           DAY           DAY_OF_WEEK
##  Min.    :2015    Min.      : 1.000    Min.      : 1.0    Min.      :1.000
##  1st Qu.:2015    1st Qu.:  3.000    1st Qu.:  8.0    1st Qu.:2.000
##  Median :2015    Median :  6.000    Median :16.0    Median :4.000
##  Mean   :2015    Mean   :  6.198    Mean   :15.7    Mean   :3.934
##  3rd Qu.:2015    3rd Qu.:  9.000    3rd Qu.:23.0    3rd Qu.:6.000
##  Max.   :2015    Max.   :12.000    Max.   :31.0    Max.   :7.000
##  AIRLINE           FLIGHT_NUMBER
##  Length:389369    Min.      :   1
##  Class :character  1st Qu.: 501
##  Mode  :character  Median :1296
##                               Mean   :1905
##                               3rd Qu.:2617
##                               Max.   :6896
```

... But for categorical data this is not very insightful, as we can see for the `AIRLINE` column

## Inspecting categorical variables

First we list all **unique elements** using in a categorical variable:

```
flights[, unique(AIRLINE)]
```

```
## [1] "AA" "US" "DL" "UA" "OO" "AS" "B6" "NK" "VX" "WN" "HA" "F9" "MQ"
```

Another valuable information for categorical variables is **how often** each category occurs:

```
flights[, table(AIRLINE)]
```

```
## AIRLINE
```

```
##   AA   AS   B6   DL   F9   HA   MQ   NK   OO   UA   US   VX   WN
## 65483 16144  8216 50343 2770 3112  368 8688 73389 54862 7374 23598 75022
```



## Row subsetting

## Row subsetting using the `i` argument

Remember the basic syntax:

```
DT[ i,  j,  by ] #
   |   |   |
   |   |   |-----> grouped by what?
   |   |-----> what to do with the columns?
   |-----> on which rows?
```

- The `i` argument allows row indexing
- `i` can be any vector of integers corresponding to
  - the row indices to select or
  - some logical vectors indicating which rows to select

## Subsetting a single row by index

If we want to see the second element of the table, we can do the following:

```
flights[2, ] # Access the 2nd row (also flights[2] or flights[i = 2])
```

```
##      YEAR MONTH DAY DAY_OF_WEEK AIRLINE FLIGHT_NUMBER TAIL_NUMBER ORIGIN_AIRPORT
## 1: 2015      1   1           4      AA           258      N3HYAA           LAX
##      DESTINATION_AIRPORT DEPARTURE_TIME AIR_TIME DISTANCE ARRIVAL_TIME
## 1:                MIA              15      258      2342           756
```

A shorter writing allows leaving out the comma:

```
flights[2]
```

```
##      YEAR MONTH DAY DAY_OF_WEEK AIRLINE FLIGHT_NUMBER TAIL_NUMBER ORIGIN_AIRPORT
## 1: 2015      1   1           4      AA           258      N3HYAA           LAX
##      DESTINATION_AIRPORT DEPARTURE_TIME AIR_TIME DISTANCE ARRIVAL_TIME
## 1:                MIA              15      258      2342           756
```

## Subsetting multiple rows by indices

For accessing **multiple consecutive** rows we can use the `start:stop` syntax:

```
flights[1:3]
```

##	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	TAIL_NUMBER	ORIGIN_AIRPORT
## 1:	2015	1	1	4	AA	2336	N3KUA	LAX
## 2:	2015	1	1	4	AA	258	N3HYA	LAX
## 3:	2015	1	1	4	US	2013	N584UW	LAX
##	DESTINATION_AIRPORT	DEPARTURE_TIME	AIR_TIME	DISTANCE	ARRIVAL_TIME			
## 1:	PBI	2	263	2330	741			
## 2:	MIA	15	258	2342	756			
## 3:	CLT	44	228	2125	753			

## Subsetting multiple rows by indices

Accessing multiple rows that are **not necessarily consecutive** can be done by creating an index vector with `c()`:

```
flights[c(3, 5)]
```

##	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	TAIL_NUMBER	ORIGIN_AIRPORT
## 1:	2015	1	1	4	US	2013	N584UW	LAX
## 2:	2015	1	1	4	AA	115	N3CTAA	LAX
##	DESTINATION_AIRPORT	DEPARTURE_TIME	AIR_TIME	DISTANCE	ARRIVAL_TIME			
## 1:	CLT	44	228	2125	753			
## 2:	MIA	103	255	2342	839			

## Subsetting rows by logical conditions

- Often, a more useful way to subset rows is using logical conditions, using for `i` a logical vector
- We can create such logical vectors using the following binary operators:
  - `==`
  - `<`
  - `>`
  - `!=`
  - `%in%`

## Subsetting rows by logical conditions with ==

For example, entries of flights operated by “AA” (American Airlines) can be extracted using:

```
flights_subset <- flights[AIRLINE == "AA"]
head(flights_subset)
```

##	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	TAIL_NUMBER	ORIGIN_AIRPORT
## 1:	2015	1	1	4	AA	2336	N3KUA	LAX
## 2:	2015	1	1	4	AA	258	N3HYA	LAX
## 3:	2015	1	1	4	AA	115	N3CTA	LAX
## 4:	2015	1	1	4	AA	2410	N3BAA	LAX
## 5:	2015	1	1	4	AA	1515	N3HMA	LAX
## 6:	2015	1	1	4	AA	1686	N4XXA	LAX
##	DESTINATION_AIRPORT			DEPARTURE_TIME	AIR_TIME	DISTANCE	ARRIVAL_TIME	
## 1:	PBI			2	263	2330	741	
## 2:	MIA			15	258	2342	756	
## 3:	MIA			103	255	2342	839	
## 4:	DFW			600	150	1235	1052	
## 5:	ORD			557	202	1744	1139	
## 6:	STL			609	183	1592	1134	

## Subsetting rows by logical conditions with %in%

We are now interested in all flights from any destination to the airports in NYC ("JFK" and "LGA"):

```
flights_subset <- flights[DESTINATION_AIRPORT %in% c("LGA", "JFK")]
tail(flights_subset)
```

##	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	TAIL_NUMBER	ORIGIN_AIRPORT
## 1:	2015	12	31	4	VX	416	N629VA	LAX
## 2:	2015	12	31	4	AA	180	N796AA	LAX
## 3:	2015	12	31	4	B6	524	N934JB	LAX
## 4:	2015	12	31	4	B6	624	N942JB	LAX
## 5:	2015	12	31	4	DL	1262	N394DL	LAX
## 6:	2015	12	31	4	B6	1124	N943JB	LAX
##	DESTINATION_AIRPORT	DEPARTURE_TIME	AIR_TIME	DISTANCE	ARRIVAL_TIME			
## 1:	JFK	1815	252	2475	201			
## 2:	JFK	1640	259	2475	18			
## 3:	JFK	1645	261	2475	18			
## 4:	JFK	2107	280	2475	513			
## 5:	JFK	2244	256	2475	625			
## 6:	JFK	2349	274	2475	748			



## Subsetting rows by logical conditions with | and &

We can concatenate multiple conditions using the logical OR | or the logical AND & operator:

```
flights_subset <- flights[AIRLINE=="AA" & DEPARTURE_TIME>600 & DEPARTURE_TIME<700]
tail(flights_subset)
```

##	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	TAIL_NUMBER	ORIGIN_AIRPORT
## 1:	2015	12	31	4	AA	700	N563UW	LAX
## 2:	2015	12	31	4	AA	169	N787AA	SFO
## 3:	2015	12	31	4	AA	1352	N7CAAA	MIA
## 4:	2015	12	31	4	AA	146	N3MKAA	LAX
## 5:	2015	12	31	4	AA	2453	N869AA	LAX
## 6:	2015	12	31	4	AA	118	N791AA	LAX
##	DESTINATION_AIRPORT			DEPARTURE_TIME	AIR_TIME	DISTANCE	ARRIVAL_TIME	
## 1:	PHL			620	252	2402	1402	
## 2:	LAX			623	54	337	740	
## 3:	LAX			651	303	2342	913	
## 4:	BOS			650	268	2611	1446	
## 5:	DFW			651	142	1235	1134	
## 6:	JFK			659	272	2475	1505	

## Column operations

## data.table environment

**Why does R correctly run code such as `flights[AIRLINE == "AA"]`?**

- Remember: `AIRLINE` is not a variable but a column of the `data.table` `flights`
- Such a call would not execute properly with a `data.frame`
- Answer: code entered inside the `[]` brackets of a `data.table` is interpreted using the `data.table` environment
  - Inside this environment, columns are seen as variables already
  - This makes the syntax very light and readable for row subsetting
  - It becomes particularly powerful for column operations

## Accessing columns but names

- Although feasible, it is not advisable to access a column by its number since
  - the ordering or number of columns can easily change.
  - Also, if you have a data set with a large number of columns (e.g. 50), how do you know which one is column 18?
- Therefore, **use the column names** to access columns for
  - preventing bugs and
  - more readability: `flights[, TAIL_NUMBER]` instead of `flights[, 7]`

## Accessing one column

```
flights[1:10, TAIL_NUMBER]      # Access column x (also DT$x or DT[j=x]).
```

```
## [1] "N3KUA" "N3HYA" "N584UW" "N547US" "N3CTAA" "N76517" "N925SW" "N719SK"  
## [9] "N435SW" "N560SW"
```

For accessing a specific cell (i.e. specific column and specific row), we can use the following syntax:

```
flights[4, TAIL_NUMBER]      # Access a specific cell.
```

```
## [1] "N547US"
```

## Accessing multiple columns

This command for accessing multiple columns would return a vector:

```
flights[1:2, c(TAIL_NUMBER, ORIGIN_AIRPORT)]
```

```
## [1] "N3KUAA" "N3HYAA" "LAX"      "LAX"
```

However, when accessing many columns, we probably want to return a `data.table` instead of a vector. For that, we need to provide R with a list, so we use `list(colA, colB)` or its simplified version `.(colA, colB)`:

```
flights[1:2, list(TAIL_NUMBER, ORIGIN_AIRPORT)]
```

```
##      TAIL_NUMBER ORIGIN_AIRPORT
## 1:      N3KUAA           LAX
## 2:      N3HYAA           LAX
```

*# Same as before.*

```
flights[1:2, .(TAIL_NUMBER, ORIGIN_AIRPORT)]
```

```
##      TAIL_NUMBER ORIGIN_AIRPORT
## 1:      N3KUAA           LAX
## 2:      N3HYAA           LAX
```

## Column operations

Since columns are seen as variables inside the `[]` environment, we can apply functions to them:

```
# Similar to mean(flights[, AIR_TIME])  
flights[, mean(AIR_TIME, na.rm=TRUE)]
```

```
## [1] 162.1379
```

```
flights[AIRLINE == "OO", mean(AIR_TIME, na.rm=TRUE)]
```

```
## [1] 68.02261
```

## Multiple column operations

- To compute operations in multiple columns, we must provide a list (unless we want the result to be a vector).

```
# Same as flights[, .(mean(AIR_TIME), median(AIR_TIME))]  
flights[, list(mean(AIR_TIME, na.rm=TRUE), median(AIR_TIME, na.rm=TRUE))]
```

```
##           V1  V2  
## 1: 162.1379 150
```

- To give meaningful names to the computations from before, we can use the following command:

```
flights[, .(mean_AIR_TIME = mean(AIR_TIME, na.rm=TRUE),  
           median_AIR_TIME = median(AIR_TIME, na.rm=TRUE))]
```

```
##      mean_AIR_TIME median_AIR_TIME  
## 1:           162.1379           150
```



## Column operations

- Any operation can be applied to the columns, just as with variables
- This code computes the average speed as the ratio of AIR\_TIME over DISTANCE for the 5 first entries of the table `flights`:

```
flights[1:5,AIR_TIME/DISTANCE]
```

```
## [1] 0.1128755 0.1101623 0.1072941 0.1224756 0.1088813
```

# Grouping

## The 'by' option

The `by` option allows executing the `j` command by groups. For example, we can use `by =` to compute the mean flight time per airline:

```
flights[, .(mean_AIRTIME = mean(AIR_TIME, na.rm=TRUE)), by = AIRLINE]
```

```
##      AIRLINE mean_AIRTIME
## 1:      AA      219.48133
## 2:      US      210.39488
## 3:      DL      207.07201
## 4:      UA      211.62008
## 5:      OO       68.02261
## 6:      AS      141.01870
## 7:      B6      309.79568
## 8:      NK      179.55828
## 9:      VX      185.36374
## 10:     WN      105.19976
## 11:     HA      307.95961
## 12:     F9      159.94041
## 13:     MQ      102.15210
```

## The 'by' option

We can also compute the mean and standard deviation of the air time of every airline:

```
flights[, .(mean_AIRTIME = mean(AIR_TIME, na.rm=TRUE),
  sd_AIR_TIME = sd(AIR_TIME, na.rm=TRUE)), by = AIRLINE]
```

##	AIRLINE	mean_AIRTIME	sd_AIR_TIME
## 1:	AA	219.48133	92.889719
## 2:	US	210.39488	105.224833
## 3:	DL	207.07201	88.908566
## 4:	UA	211.62008	94.832456
## 5:	OO	68.02261	41.065036
## 6:	AS	141.01870	51.806424
## 7:	B6	309.79568	28.457740
## 8:	NK	179.55828	78.194706
## 9:	VX	185.36374	113.504572
## 10:	WN	105.19976	69.257334
## 11:	HA	307.95961	23.905491
## 12:	F9	159.94041	61.412379
## 13:	MQ	102.15210	8.531046

## Remark on the `data.table` syntax

- Although we could write `flights[i = 5, j = AIRLINE]`, we usually omit the `i =` and `j =` from the syntax, and write `flights[5, AIRLINE]` instead.
- However, for clarity we usually include the `by =` in the syntax.

## Counting occurrences with `.N`

## Counting occurrences with `.N`

The `.N` is a special in-built variable that counts the number observations within a table. Evaluating `.N` alone is equal to `nrow()` of a table:

```
flights[, .N]
```

```
## [1] 389369
```

```
nrow(flights)
```

```
## [1] 389369
```

## Powerful statements using all three elements i, j and by

Remember the data.table definition: “Take **DT**, subset rows using **i**, then select or calculate **j**, grouped by **by**”

For example, we can, for each airline, get the number of flights arriving to the airport JFK:

```
flights[DESTINATION_AIRPORT == "JFK", .N, by = 'AIRLINE']
```

```
##      AIRLINE      N
## 1:         B6 2488
## 2:         DL 2546
## 3:         AA 3804
## 4:         VX 1652
## 5:         UA 1525
```



## Extending tables

## Creating new columns (the := command)

The := operator updates the data.table inplace, so writing `DT <- DT[, ... := ...]` is redundant.

This operator changes the input by *reference*. No copy of the object is made, which makes the operation faster and less memory-consuming.

As an example, we can add a new column called `SPEED` (in miles per hour) whose value is the `DISTANCE` divided by `AIR_TIME` times 60:

```
flights[, SPEED := DISTANCE / AIR_TIME * 60]
head(flights)
```

##	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	TAIL_NUMBER	ORIGIN_AIRPORT
## 1:	2015	1	1	4	AA	2336	N3KUAA	LAX
## 2:	2015	1	1	4	AA	258	N3HYAA	LAX
## 3:	2015	1	1	4	US	2013	N584UW	LAX
## 4:	2015	1	1	4	DL	1434	N547US	LAX
## 5:	2015	1	1	4	AA	115	N3CTAA	LAX
## 6:	2015	1	1	4	UA	1545	N76517	LAX
##	DESTINATION_AIRPORT			DEPARTURE_TIME	AIR_TIME	DISTANCE	ARRIVAL_TIME	SPEED
## 1:	PBI			2	263	2330	741	531.5589
## 2:	MIA			15	258	2342	756	544.6512
## 3:	CLT			44	228	2125	753	559.2105
## 4:	MSP			35	188	1535	605	489.8936
## 5:	MIA			103	255	2342	839	551.0588
## 6:	IAH			112	156	1379	607	530.3846

## Creating and using new columns (the := command)

Having computed a new column using the := operator, we can use it for further analyses.

For instance, we can compute the average speed, air time and distance for each airline:

```
flights[, .(mean_AIR_TIME = mean(AIR_TIME, na.rm=TRUE),
      mean_SPEED = mean(SPEED, na.rm=TRUE),
      mean_DISTANCE = mean(DISTANCE, na.rm=TRUE)
    ), by=AIRLINE]
```

```
##      AIRLINE mean_AIR_TIME mean_SPEED mean_DISTANCE
##  1:      AA      219.48133      461.2839      1739.2331
##  2:      US      210.39488      452.1641      1658.2581
##  3:      DL      207.07201      466.0330      1656.2165
##  4:      UA      211.62008      464.2928      1693.5504
##  5:      OO       68.02261      349.5549       437.2337
##  6:      AS      141.01870      439.0120      1040.0340
##  7:      B6      309.79568      484.8242      2486.1489
##  8:      NK      179.55828      450.0221      1402.1591
##  9:      VX      185.36374      433.0870      1432.5384
## 10:      WN      105.19976      409.3803       760.2593
## 11:      HA      307.95961      497.3118      2537.8107
## 12:      F9      159.94041      461.0684      1235.6664
## 13:      MQ      102.15210      435.5580       737.0000
```

## Removing columns

Additionally we can use the `:=` operator to remove columns.

If we for example observe that tail numbers are not important for our analysis we can remove them with the following statement:

```
flights[, TAIL_NUMBER := NULL]
head(flights)
```

```
##      YEAR MONTH DAY DAY_OF_WEEK AIRLINE FLIGHT_NUMBER ORIGIN_AIRPORT
## 1: 2015      1  1           4      AA          2336           LAX
## 2: 2015      1  1           4      AA           258           LAX
## 3: 2015      1  1           4      US          2013           LAX
## 4: 2015      1  1           4      DL          1434           LAX
## 5: 2015      1  1           4      AA           115           LAX
## 6: 2015      1  1           4      UA          1545           LAX
##      DESTINATION_AIRPORT DEPARTURE_TIME AIR_TIME DISTANCE ARRIVAL_TIME   SPEED
## 1:                    PBI              2     263    2330        741 531.5589
## 2:                    MIA             15     258    2342        756 544.6512
## 3:                    CLT             44     228    2125        753 559.2105
## 4:                    MSP             35     188    1535        605 489.8936
## 5:                    MIA            103     255    2342        839 551.0588
## 6:                    IAH            112     156    1379        607 530.3846
```

# Copying tables

What do we mean when we say that `data.table` modifies columns *by reference*?

It means that no new copy of the object is made in the memory, unless we actually create one using `copy()`.

```
or_dt <- data.table(a = 1:10, b = 11:20)
# No new object is created, both new_dt and or_dt point to the same memory chunk.
new_dt <- or_dt
new_dt[, ab := a*b]
colnames(or_dt) # or_dt was also affected by changes in new_dt
```

```
## [1] "a" "b" "ab"
```

```
or_dt <- data.table(a = 1:10, b = 11:20)
copy_dt <- copy(or_dt) # By creating a copy, we have 2 objects in memory
copy_dt[, ab := a*b]
colnames(or_dt) # Changes in the copy don't affect the original
```

```
## [1] "a" "b"
```

# Data.table exercises

## Summary

By now, you should be able to answer the following questions:

- How to subset by rows or columns? Remember: `DT[i, j, by]`.
- How to add columns?
- How to make operations with different columns?

## Data.table resources

The help page for `data.table`.

<https://cran.r-project.org/web/packages/data.table/>

<https://s3.amazonaws.com/..assets.datacamp.com/img/blog/data+table+cheat+sheet.pdf>

<http://r4ds.had.co.nz/relational-data.html>

<http://adv-r.had.co.nz/Environments.html>



# Rmarkdown

# Creating reproducible reports

- This is an R Markdown presentation. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.
- Simply go to File → New File → R Markdown
- Select PDF and you get a template.
- All the commands that you may need can be found on this cheatsheet: <https://raw.githubusercontent.com/rstudio/cheatsheets/master/rmarkdown-2.0.pdf>.
- When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document.

# Recap

# In a nutshell

Today we learned:

- The basics of R and of the Rstudio interface
- How to manipulate the basic R data types
- How to load datasets and do operations on them with `data.table`
- How to make reproducible reports

# Reading

## Reading

You can read more details on the subjects that we discussed today on **Chapter 1: R basics** and **Chapter 2: Data wrangling**, as well as the **A: Importing data** and **B: R programming** of the **Appendix**, on <https://gagneurlab.github.io/dataviz/>.

# Cheatsheets

# Cheatsheets

You can find the whole book of our Data Analysis and Visualization in R lecture here:  
[<https://gagneurlab.github.io/dataviz/>]

A cheatsheet for R studio can be obtained here:  
[<https://raw.githubusercontent.com/rstudio/cheatsheets/master/rstudio-ide.pdf>]

A cheatsheet for R basics can be obtained here:  
[<https://www.rstudio.com/wp-content/uploads/2016/10/r-cheat-sheet-3.pdf>]

A cheatsheet for simple data.table manipulations can be obtained here: [<https://datacamp-community-prod.s3.amazonaws.com/6fdf799f-76ba-45b1-b8d8-39c4d4211c31>]

A cheatsheet for advanced data.table manipulations can be obtained here:  
[<https://raw.githubusercontent.com/rstudio/cheatsheets/master/datatable.pdf>]

A cheatsheet on how to create Rmarkdowns can be obtained here:  
[<https://raw.githubusercontent.com/rstudio/cheatsheets/master/rmarkdown-2.0.pdf>]